



S.B. JAIN INSTITUTE OF TECHNOLOGY MANAGEMENT & RESEARCH, NAGPUR

Practical 05

Aim: Write a program to implement Shortest Job First (SJF) Preemptive Scheduling for three processes and calculate the total context switches and average waiting time. The processes have burst times 10ns, 20ns, and 30ns, arriving at 0ns, 2ns, and 6ns, respectively.

Name: Sayali Shamrao Misal

USN: CM25D007

Semester/Year: IV/II

Academic Session: 2025-2026

Date of Performance:

Date of Submission:

❖ **Aim:** Write a program to implement Shortest Job First (SJF) Preemptive Scheduling for three processes and calculate the total context switches and average waiting time. The processes have burst times 10ns, 20ns, and 30ns, arriving at 0ns, 2ns, and 6ns, respectively.

❖ **Objectives:**

Understand SJF Preemptive Scheduling: Implement the **Shortest Job First (SJF) Preemptive Scheduling** algorithm to manage CPU execution efficiently.

Calculate Context Switches: Determine the total number of context switches required for the given set of processes.

Evaluate Waiting Time: Compute the **average waiting time** for all processes before getting CPU execution.

❖ **Requirements:**

✓ **Hardware Requirements:**

- Processor: Minimum 1 GHz
- RAM: 512 MB or higher
- Storage: 100 MB free space

✓ **Software Requirements:**

- Operating System: Linux/Unix-based
- Shell: Bash 4.0 or higher
- Text Editor: Nano, Vim, or any preferred editor

❖ **Theory:**

CPU Scheduling in Operating Systems

Introduction

Scheduling is the method by which processes are given access to the CPU. Efficient scheduling is essential for optimal system performance and user experience. There are two primary types of CPU scheduling: **Preemptive Scheduling** and **Non-Preemptive Scheduling**.

Understanding the differences between these scheduling types helps in designing and choosing the right scheduling algorithms for different operating systems.

1. Preemptive Scheduling

In **Preemptive Scheduling**, the operating system can interrupt or preempt a running process to allocate CPU time to another process, typically based on priority or time-sharing policies.

Algorithms Based on Preemptive Scheduling:

- **Round Robin (RR)**
- **Shortest Remaining Time First (SRTF)**
- **Priority Scheduling (Preemptive version)**

Example:

In the following case, **P2** is preempted at time 1 due to the arrival of a higher-priority process.

Advantages of Preemptive Scheduling:

- ✓ Prevents a process from monopolizing the CPU, improving system reliability.
- ✓ Enhances **average response time**, making it beneficial for multi-programming environments.

2. Non-Preemptive Scheduling

In **Non-Preemptive Scheduling**, a running process cannot be interrupted by the operating system. It continues executing until it **terminates** or **enters a waiting state** voluntarily.

Algorithms Based on Non-Preemptive Scheduling:

- **First Come First Serve (FCFS)**
- **Shortest Job First (SJF - Non-Preemptive)**
- **Priority Scheduling (Non-Preemptive version)**

Example:

Below is a **Gantt Chart** based on the **FCFS algorithm**, where each process executes fully before the next one starts.

Advantages of Non-Preemptive Scheduling:

- ✓ **Easy to implement** in an operating system (used in older versions like Windows 3.11 and early macOS).
- ✓ **Minimal scheduling overhead** due to fewer context switches.

3. Preemptive Scheduling

In **Preemptive Scheduling**, the operating system can interrupt or preempt a running process to allocate CPU time to another process, typically based .

Algorithms Based on Preemptive Scheduling:

- **Round Robin (RR)**
- **Shortest Remaining Time First (SRTF)**
- **Priority Scheduling (Preemptive version)**

Example:

In the following case, **P2 is preempted at time 1** due to the arrival of a higher-priority process.

Advantages of Preemptive Scheduling:

- ✓ Prevents a process from monopolizing the CPU, improving system reliability.
- ✓ Enhances **average response time**, making it beneficial for multi-programming environments.
- ✓ Used in modern operating systems like **Windows, Linux, and macOS**.

Disadvantages of Preemptive Scheduling:

4. Non-Preemptive Scheduling

In **Non-Preemptive Scheduling**, a running process cannot be interrupted by the operating system. It continues executing until it **terminates** or **enters a waiting state** voluntarily.

Algorithms Based on Non-Preemptive Scheduling:

- **First Come First Serve (FCFS)**
- **Shortest Job First (SJF - Non-Preemptive)**
- **Priority Scheduling (Non-Preemptive version)**

Example:

Below is a **Gantt Chart** based on the **FCFS algorithm**, where each process executes fully before the next one starts.

Advantages of Non-Preemptive Scheduling:

- ✓ **Easy to implement** in an operating system (used in older versions like Windows 3.11 and early macOS).
- ✓ **Minimal scheduling overhead** due to fewer context switches.

5. Differences Between Preemptive and Non-Preemptive Scheduling

Parameter	Preemptive Scheduling	Non-Preemptive Scheduling
Basic Concept	CPU time is allocated for a limited time .	CPU is held until process terminates or enters waiting state.

Starvation	Frequent high-priority processes may starve low-priority ones.	A long-running process can starve later-arriving shorter processes.
Overhead	Higher overhead due to frequent context switching .	Minimal overhead.
Flexibility	More flexible (critical processes get priority).	Rigid scheduling approach.
Response Time	Faster response time.	Slower response time.
Process Control	OS has more control over scheduling.	OS has less control over scheduling.
Concurrency Issues	Higher , as processes may be preempted during shared resource access.	Lower , as processes run to completion.
Examples	Round Robin, SRTF.	FCFS, Non-Preemptive SJF.

6. Frequently Asked Questions (FAQs)

a. How is priority determined in Preemptive scheduling?

Ans: Preemptive scheduling systems assign priority based on **task importance, deadlines, or urgency**. Higher-priority tasks execute before lower-priority ones.

b. What happens in non-preemptive scheduling if a process does not yield the CPU?

Ans: If a process does not voluntarily yield the CPU, it can lead to **starvation or deadlock**, where other tasks are unable to execute.

c. Which scheduling method is better for real-time systems?

Ans: Preemptive scheduling is better for **real-time systems**, as it allows high-priority tasks to execute immediately.

Conclusion: Preemptive scheduling offers better responsiveness but adds complexity, while non-preemptive scheduling is simpler but may cause inefficiencies. The choice depends on system needs, with preemptive suited for multitasking and non-preemptive for low-overhead scenarios

- **Discussion Questions:**

1. **What is the key difference between preemptive and non-preemptive scheduling?**

2. Why does preemptive scheduling require context switching?
3. Which CPU scheduling algorithm is most suitable for real-time systems and why?
4. What is starvation in CPU scheduling, and how can it be prevented?
5. Why is the Round Robin scheduling algorithm preferred in time-sharing systems?

❖ References:

<https://www.geeksforgeeks.org/preemptive-and-non-preemptive-scheduling/>

Date: _____ / _____ /2026

Signature

Course Coordinator
B.Tech CSE(AIML)
Sem: 4 / 2025-26

```

sayali@DESKTOP-T8KMP91:~$ nano sjf.c
sayali@DESKTOP-T8KMP91:~$ gcc sjf.c -o sjf
sayali@DESKTOP-T8KMP91:~$ ./sjf

Process Arrival Burst Waiting
P1      0        10      0
P2      2        20      8
P3      6        30     24

Average Waiting Time = 10.67 ns
Total Context Switches = 2
sayali@DESKTOP-T8KMP91:~$ nano sjf.c
sayali@DESKTOP-T8KMP91:~$ gcc sjf.c -o sjf
sayali@DESKTOP-T8KMP91:~$ ./sjf

Gantt Chart:
| P1 | P2 |
| P2 | P3 |
| P3 |
| P3 |
0   1   2   3   4   5   6   7   8   9   10  11  12  13  14  15
16  17  18  19  20  21  22  23  24  25  26  27  28  29
30  31  32  33  34  35  36  37  38  39  40  41  42  43   4
4   45  46  47  48  49  50  51  52  53  54  55  56  57   58
59  60

```

CODE:

```

#include <stdio.h>
#include <limits.h>

int main() {
    int n = 3;
    int at[3] = {0, 2, 6}; // Arrival times
    int bt[3] = {10, 20, 30}; // Burst times
    int rt[3]; // Remaining times
    int wt[3] = {0}; // Waiting times
    int ct[3] = {0}; // Completion times

    int time = 0, completed = 0;
    int min, shortest;
    int finish_time;
    int check;
    int prev = -1;
    int context_switch = 0;

    // Copy burst time into remaining time
    for (int i = 0; i < n; i++) {
        rt[i] = bt[i];
    }

    while (completed != n) {
        min = INT_MAX;
        check = 0;

        for (int i = 0; i < n; i++) {
            if ((at[i] <= time) && (rt[i] < min) && rt[i] > 0) {
                min = rt[i];
                shortest = i;
                check = 1;
            }
        }

        if (check == 0) {
            time++;
            continue;
        }

        // Count context switch
        if (prev != -1 && prev != shortest) {
            context_switch++;
        }
        prev = shortest;

        // Execute process for 1 time unit
        rt[shortest]--;
        time++;

        // If process finishes
        if (rt[shortest] == 0) {
            completed++;
            finish_time = time;

            ct[shortest] = finish_time;
            wt[shortest] = ct[shortest] - at[shortest] - bt[shortest];

            if (wt[shortest] < 0)
                wt[shortest] = 0;
        }
    }

    // Calculate average waiting time
    float total_wt = 0;

    printf("\nProcess\tArrival\tBurst\tWaiting\n");
    for (int i = 0; i < n; i++) {
        total_wt += wt[i];
        printf("P%d\t%d\t%d\t%d\t%d\n", i + 1,
               at[i], bt[i], wt[i]);
    }

    printf("\nAverage Waiting Time = %.2f ns",
           total_wt / n);
    printf("\nTotal Context Switches = %d\n", context_switch);

    return 0;
}

```